

# Assignment 6 - Cython

October 27, 2024

## Numerical Integration Using the Trapezoidal Rule: Python vs. Cython vs. NumPy

### 1. Introduction

Numerical integration approximates definite integrals of functions lacking closed-form solutions. The trapezoidal rule, favored for its simplicity and accuracy, serves as an ideal method to compare implementations in Python, Cython, and NumPy, focusing on accuracy and performance enhancements.

### 2. Implementations

#### a. Python (py\_trapz.py)

A straightforward approach calculates trapezoid widths and sums function values iteratively.

```
def py_trapz(f, a, b, n):
    h = (b - a) / n
    integral = 0.5 * (f(a) + f(b))
    for i in range(1, n):
        integral += f(a + i * h)
    return integral * h
```

#### b. Cython (cy\_trapz.pyx)

Cython optimizes performance by adding static types and compiler directives, compiling to C for speed.

```
cimport cython

@cython.boundscheck(False)
@cython.wraparound(False)
cpdef double cy_trapz(object f, double a, double b, long n):
    cdef double h = (b - a) / n
    cdef double integral = 0.5 * (f(a) + f(b))
    cdef long i
    for i in range(1, n):
        integral += f(a + i * h)
    return integral * h
```

### 3. Accuracy

All implementations achieved high accuracy across test cases, closely matching exact integrals with minimal errors ( $10^{-10}$  to  $10^{-12}$ ), validating their reliability.

## 4. Performance

- **Small Integrations ( $n = 1,000$ ):**
  - **Cython** was fastest due to optimized loops.
  - **NumPy** slightly slower from array overhead.
  - **Python** was the slowest.
- **Large Integrations ( $n = 10,000,000$ ):**
  - **NumPy** excelled with its optimized backend.
  - **Cython** remained faster than Python.
  - **Python** lagged significantly.

## 5. Comparison Against NumPy

Cython outperforms Python in loop-intensive tasks and smaller datasets by optimizing execution. However, NumPy surpasses Cython in large-scale integrations due to its highly optimized C/Fortran backend and vectorized operations.

## 6. Challenges and Optimizations in Cython

Implementing Cython involved:

- **Learning Cython Syntax:** Differentiating `def`, `cdef`, and `cpdef` for optimal performance.
- **Compiler Directives:** Disabling bounds checking and wraparound to enhance speed.
- **Integration Issues:** Ensuring seamless interaction with Python functions and proper benchmarking using `timeit`.

## 7. Lessons Learned

- **Cython's Efficiency:** Significantly boosts performance for computation-heavy tasks.
- **Static Typing:** Reduces overhead and speeds up execution.
- **Balanced Accessibility:** `cpdef` functions offer both Python usability and C-level speed.
- **Accurate Benchmarking:** Essential for true performance assessment, avoiding scope and import issues.

## 8. Conclusion

The trapezoidal rule implementations highlight distinct strengths:

- **Cython** is ideal for small to medium-sized, loop-driven integrations, offering substantial speed improvements over pure Python.
- **NumPy** is superior for large-scale numerical operations, leveraging its optimized backend for maximum performance.

Choosing between Cython and NumPy depends on the specific computational requirements, with Cython enhancing Python's capabilities for intensive tasks and NumPy providing unmatched efficiency for large datasets.